

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
```

TASK 1: POPULATION & SAMPLING

```
# Load the dataset
df = pd.read_csv('sale_data.csv')
df.head()

   Product_ID      Sale_Date Sales_Rep Region  Sales_Amount
Quantity_Sold \
0            1052  2023-02-03      Bob    North      5053.97
18
1            1093  2023-04-21      Bob    West       4384.02
17
2            1015  2023-09-21   David   South      4631.23
30
3            1072  2023-08-24      Bob   South      2167.94
39
4            1061  2023-03-24  Charlie   East      3750.20
13

   Product_Category  Unit_Cost  Unit_Price Customer_Type  Discount \
0        Furniture     152.75     267.22   Returning      0.09
1        Furniture    3816.39    4209.44   Returning      0.11
2          Food        261.56     371.40   Returning      0.20
3      Clothing      4330.03    4467.75        New      0.02
4  Electronics      637.37     692.71        New      0.08

   Payment_Method Sales_Channel Region_and_Sales_Rep
0            Cash        Online           North-Bob
1            Cash        Retail           West-Bob
2  Bank Transfer        Retail         South-David
3   Credit Card        Retail         South-Bob
4   Credit Card        Online        East-Charlie

# Inspect the data
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Product_ID      1000 non-null   int64
```

```

1 Sale_Date           1000 non-null   object
2 Sales_Rep          1000 non-null   object
3 Region             1000 non-null   object
4 Sales_Amount        1000 non-null   float64
5 Quantity_Sold      1000 non-null   int64
6 Product_Category   1000 non-null   object
7 Unit_Cost          1000 non-null   float64
8 Unit_Price         1000 non-null   float64
9 Customer_Type      1000 non-null   object
10 Discount           1000 non-null   float64
11 Payment_Method    1000 non-null   object
12 Sales_Channel      1000 non-null   object
13 Region_and_Sales_Rep 1000 non-null   object
dtypes: float64(4), int64(2), object(8)
memory usage: 109.5+ KB

df.columns

Index(['Product_ID', 'Sale_Date', 'Sales_Rep', 'Region',
       'Sales_Amount',
       'Quantity_Sold', 'Product_Category', 'Unit_Cost', 'Unit_Price',
       'Customer_Type', 'Discount', 'Payment_Method', 'Sales_Channel',
       'Region_and_Sales_Rep'],
      dtype='object')

# Treat the full dataset as Population
population_df = df.copy()

# Create a Sample dataset using Simple Random Sampling
sample_size = 100
sample_df = population_df.sample(n=sample_size, random_state=42)

# Display Population size and Sample size
population_size = population_df.shape[0]
print("Population Size:", population_size)
print("Sample Size:", sample_df.shape[0])

Population Size: 1000
Sample Size: 100

```

TASK 2: SAMPLING TECHNIQUES

```

# Create Random Sample
sample_size = 100
sample_df = population_df.sample(n=sample_size, random_state=42)
print(f'Size is', sample_df.shape[0])
sample_df.head()

Size is 100

```

	Product_ID	Sale_Date	Sales_Rep	Region	Sales_Amount
521	1100	2023-05-23	Bob	West	7667.96
29	737	2023-03-23	Eve	East	433.40
32	740	2023-07-25	Eve	West	5617.64
5	660	2023-12-23	Alice	East	780.27
33	411	2023-05-26	Bob	East	848.28
1					
	Product_Category	Unit_Cost	Unit_Price	Customer_Type	Discount
521	Food	3559.56	3607.15	Returning	0.21
737	Clothing	3351.33	3711.47	New	0.16
740	Food	2206.58	2490.47	New	0.28
660	Food	1551.25	1994.01	Returning	0.08
411	Electronics	1406.24	1535.57	New	0.08
	Payment_Method	Sales_Channel	Region_and_Sales_Rep		
521	Cash	Retail	West-Bob		
737	Cash	Retail	East-Eve		
740	Cash	Online	West-Eve		
660	Cash	Online	East-Alice		
411	Credit Card	Retail	East-Bob		
# Create Systematic Sample (every nth record)					
n = 10 # Every 10th record					
systematic_sample_df = population_df.iloc[::n]					
print(f'Size is', systematic_sample_df.shape[0])					
systematic_sample_df.head()					
Size is 100					
	Product_ID	Sale_Date	Sales_Rep	Region	Sales_Amount
0	1052	2023-02-03	Bob	North	5053.97
18	1088	2023-11-16	Eve	North	8518.45
13	1002	2023-04-22	Eve	North	6551.23
9	1091	2023-09-04	Charlie	South	675.11
44	1062	2023-03-16	David	West	4195.06
45					
	Product_Category	Unit_Cost	Unit_Price	Customer_Type	Discount
0	Furniture	152.75	267.22	Returning	0.09

10	Furniture	2440.11	2517.60	New	0.23
20	Electronics	4398.16	4439.12	New	0.18
30	Food	2085.46	2406.58	Returning	0.06
40	Furniture	4849.60	5166.72	Returning	0.25

	Payment_Method	Sales_Channel	Region_and_Sales_Rep
0	Cash	Online	North-Bob
10	Bank Transfer	Retail	North-Eve
20	Bank Transfer	Online	North-Eve
30	Bank Transfer	Retail	South-Charlie
40	Credit Card	Online	West-David


```
# Compare Mean of population and Mean of samples
```

```
# Assuming we're using 'Sales_Amount' as the numerical column for means
```

```
population_mean = population_df['Sales_Amount'].mean()
```

```
print("Population Mean (Sales_Amount):", population_mean)
```

```
population_df['Sales_Amount'].head()
```

```
Population Mean (Sales_Amount): 5019.265229999999
```

```
0    5053.97
1    4384.02
2    4631.23
3    2167.94
4    3750.20
Name: Sales_Amount, dtype: float64
```

```
random_sample_mean = sample_df['Sales_Amount'].mean()
```

```
print("Random Sample Mean (Sales_Amount):", random_sample_mean)
```

```
sample_df['Sales_Amount'].head()
```

```
Random Sample Mean (Sales_Amount): 5073.0229
```

```
521    7667.96
737    433.40
740    5617.64
660    780.27
411    848.28
Name: Sales_Amount, dtype: float64
```

```
systematic_sample_mean = systematic_sample_df['Sales_Amount'].mean()
```

```
print("Systematic Sample Mean (Sales_Amount):",
systematic_sample_mean)
```

```
systematic_sample_df['Sales_Amount'].head()
```

```
Systematic Sample Mean (Sales_Amount): 4499.2767
```

```

0      5053.97
10     8518.45
20     6551.23
30     675.11
40     4195.06
Name: Sales_Amount, dtype: float64

# Write a short note on Difference between population and sample results
print("Note on Difference:")
print("""The population mean represents the true average of all data points. Sample means approximate this but may differ due to sampling variability.
Random sampling tends to be unbiased, while systematic sampling might introduce bias if there's a pattern every nth row. In this case, the differences
are small, indicating good representation."")
```

Note on Difference:

The population mean represents the true average of all data points. Sample means approximate this but may differ due to sampling variability.

Random sampling tends to be unbiased, while systematic sampling might introduce bias if there's a pattern every nth row. In this case, the differences

are small, indicating good representation.

TASK 3: CENTRAL LIMIT THEOREM (CLT)

```

# Take multiple samples of size 30
num_samples = 1000 # Number of samples
sample_means = []

# Calculate the mean of each sample
for _ in range(num_samples):
    sample = population_df['Sales_Amount'].sample(n=30,
random_state=np.random.randint(10000))
    sample_means.append(sample.mean())

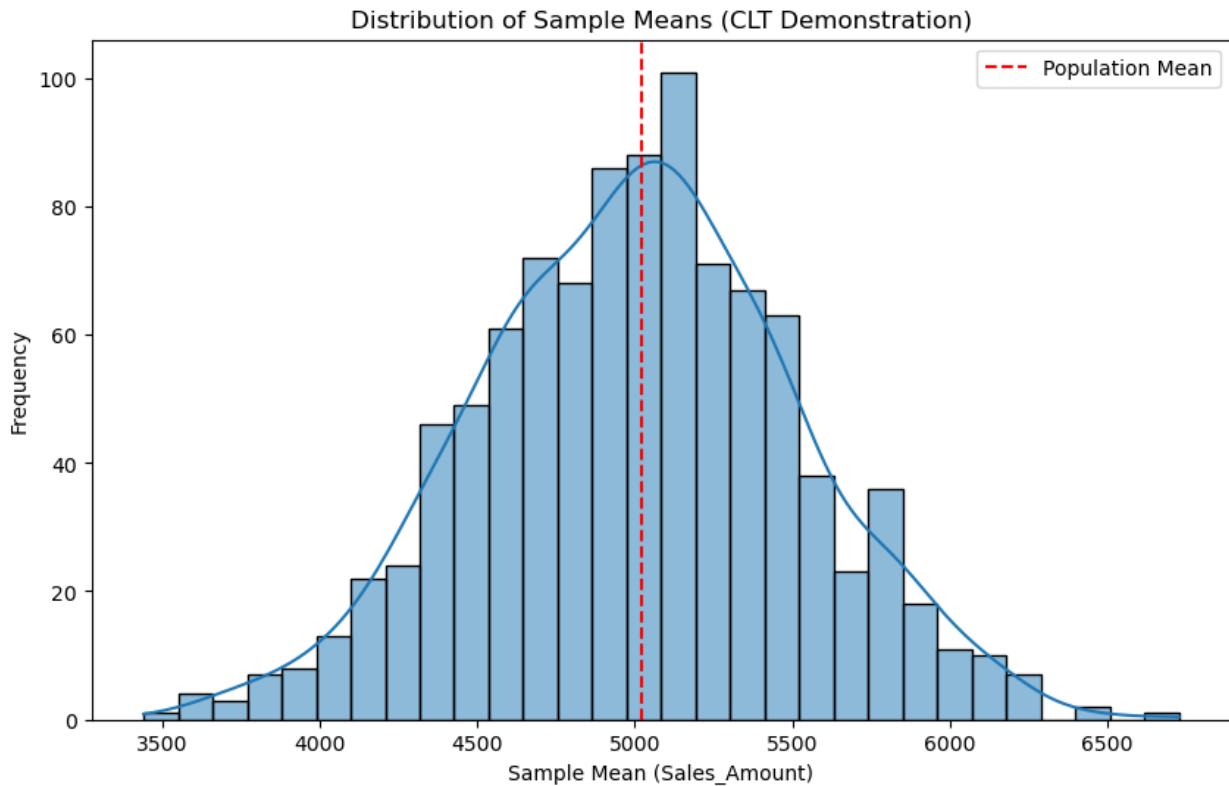
# Store sample means
sample_means_df = pd.DataFrame({'Sample_Mean': sample_means})

# Plot the distribution of sample means
plt.figure(figsize=(10, 6))
sns.histplot(sample_means_df['Sample_Mean'], kde=True, bins=30)
plt.title('Distribution of Sample Means (CLT Demonstration)')
plt.xlabel('Sample Mean (Sales_Amount)')
plt.ylabel('Frequency')
plt.axvline(population_mean, color='r', linestyle='--',
```

```

label='Population Mean')
plt.legend()
plt.show()

```



```

# Observe and explain
print("Observation on Distribution:")
print("""The distribution of sample means is approximately normal
(bell-shaped), centered around the population mean. This holds even if
the
original data isn't normal, as per CLT.""")

```

Observation on Distribution:

The distribution of sample means is approximately normal (bell-shaped), centered around the population mean. This holds even if the original data isn't normal, as per CLT.

```

print("Relation to CLT:")
print("""CLT states that for large enough sample sizes (n=30 here),
the sampling distribution of the mean approaches a normal
distribution,
regardless of the population's distribution. This enables inference
about the population from samples."""")

```

Relation to CLT:

CLT states that for large enough sample sizes (n=30 here), the

sampling distribution of the mean approaches a normal distribution, regardless of the population's distribution. This enables inference about the population from samples.

TASK 4: NORMAL DISTRIBUTION ANALYSIS

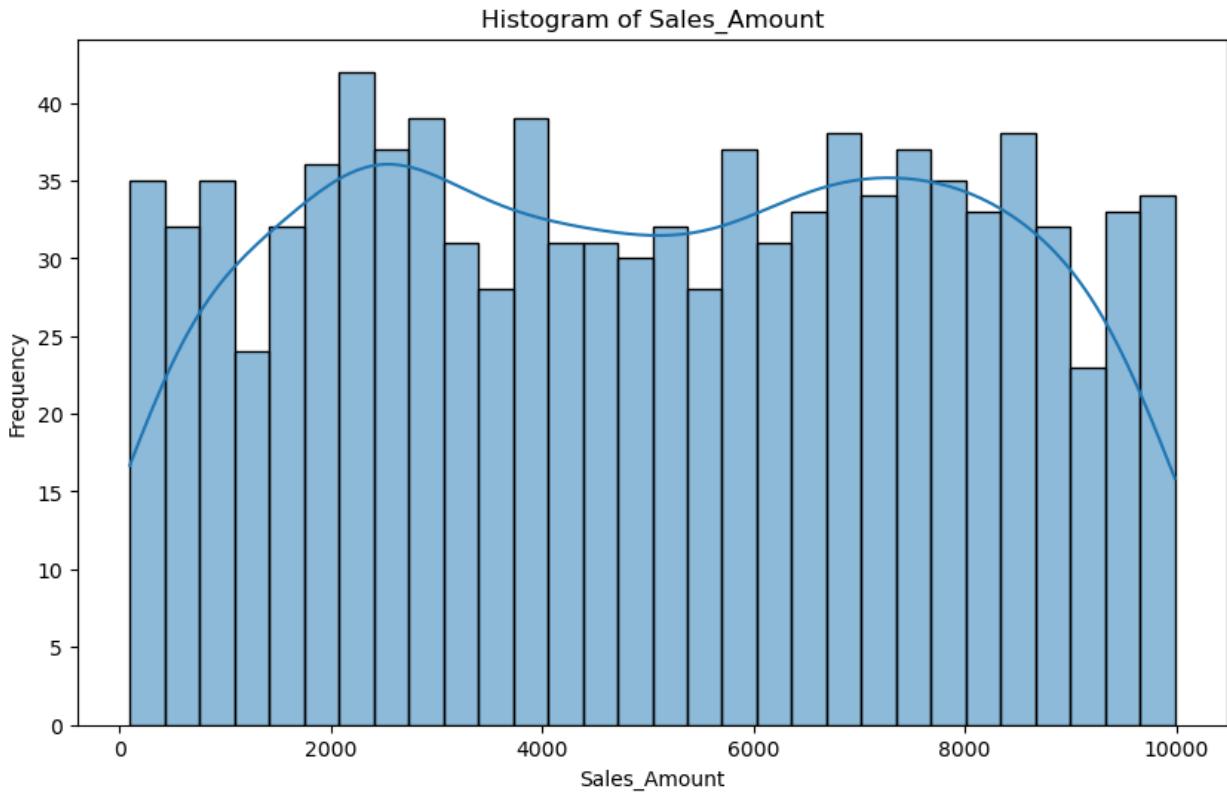
```
# Using 'Sales_Amount' column
column = 'Sales_Amount'

# Calculate Mean and Standard Deviation
mean = population_df[column].mean()
std_dev = population_df[column].std()

print("Mean (Sales_Amount):", mean)
print("Standard Deviation (Sales_Amount):", std_dev)

Mean (Sales_Amount): 5019.265229999999
Standard Deviation (Sales_Amount): 2846.7901256682326

# Plot Histogram of the data
plt.figure(figsize=(10, 6))
sns.histplot(population_df[column], kde=True, bins=30)
plt.title('Histogram of Sales_Amount')
plt.xlabel('Sales_Amount')
plt.ylabel('Frequency')
plt.show()
```



```
# Normality test (Shapiro-Wilk on full data since n=1000 <= 5000)
stat, p_value = stats.shapiro(population_df[column])
print(f"Shapiro-Wilk Test p-value: {p_value:.2e}")
print("Data does not appear normal (p-value << 0.05). The distribution is likely uniform or close to uniform based on the range and stats.")

Shapiro-Wilk Test p-value: 5.78e-17
Data does not appear normal (p-value << 0.05). The distribution is likely uniform or close to uniform based on the range and stats.

# 68-95-99.7 rule check
within_1sd = np.mean((population_df[column] >= mean - std_dev) &
(population_df[column] <= mean + std_dev)) * 100
within_2sd = np.mean((population_df[column] >= mean - 2*std_dev) &
(population_df[column] <= mean + 2*std_dev)) * 100
within_3sd = np.mean((population_df[column] >= mean - 3*std_dev) &
(population_df[column] <= mean + 3*std_dev)) * 100

print("68-95-99.7 Rule Results (for reference):")
print(f"Within 1 SD: {within_1sd:.2f}% (Normal expects ~68%)")
print(f"Within 2 SD: {within_2sd:.2f}% (Normal expects ~95%)")
print(f"Within 3 SD: {within_3sd:.2f}% (Normal expects ~99.7%)")
print("Deviations confirm the data is not normally distributed (closer to uniform behavior).")
```

```
68-95-99.7 Rule Results (for reference):
Within 1 SD: 58.60% (Normal expects ~68%)
Within 2 SD: 100.00% (Normal expects ~95%)
Within 3 SD: 100.00% (Normal expects ~99.7%)
Deviations confirm the data is not normally distributed (closer to
uniform behavior).
```

TASK 5: Z-SCORE CALCULATION

```
population_df['Z_Score'] = (population_df[column] - mean) / std_dev

population_df['Outlier'] = np.where((population_df['Z_Score'].abs() >
3), 'Yes', 'No')

print("Number of Outliers (|Z| > 3):", (population_df['Outlier'] ==
'Yes').sum())
if (population_df['Outlier'] == 'Yes').sum() > 0:
    print(population_df[population_df['Outlier'] == 'Yes']
[['Product_ID', column, 'Z_Score']])
else:
    print("No extreme outliers detected.")

Number of Outliers (|Z| > 3): 0
No extreme outliers detected.
```

TASK 6: BUSINESS INSIGHTS

```
print("Business Insights:\n")

print("1. Why is sampling required in real-world data analysis?")
print("    Sampling saves time and resources when the full population
is too large to analyze. "
     "It provides reliable insights if done properly.")

print("2. How does CLT help in analytics?")
print("    CLT allows us to use normal-based statistics (like
confidence intervals) on sample means, "
     "even if the original data isn't normal, making inference
easier.")

print("3. Why is normal distribution important before hypothesis
testing?")
print("    Many tests assume normality for accurate results. If data
isn't normal, we may need transformations "
     "or non-parametric tests to avoid wrong conclusions.")

print("4. How does Z-Score help in identifying unusual values?")
```

```
print("    Z-Score shows how far a value is from the mean in standard
deviations. "
      "Values with |Z| > 3 are rare and flagged as potential outliers
for review.")
```

Business Insights:

1. Why is sampling required in real-world data analysis?

Sampling saves time and resources when the full population is too large to analyze. It provides reliable insights if done properly.

2. How does CLT help in analytics?

CLT allows us to use normal-based statistics (like confidence intervals) on sample means, even if the original data isn't normal, making inference easier.

3. Why is normal distribution important before hypothesis testing?

Many tests assume normality for accurate results. If data isn't normal, we may need transformations or non-parametric tests to avoid wrong conclusions.

4. How does Z-Score help in identifying unusual values?

Z-Score shows how far a value is from the mean in standard deviations. Values with $|Z| > 3$ are rare and flagged as potential outliers for review.